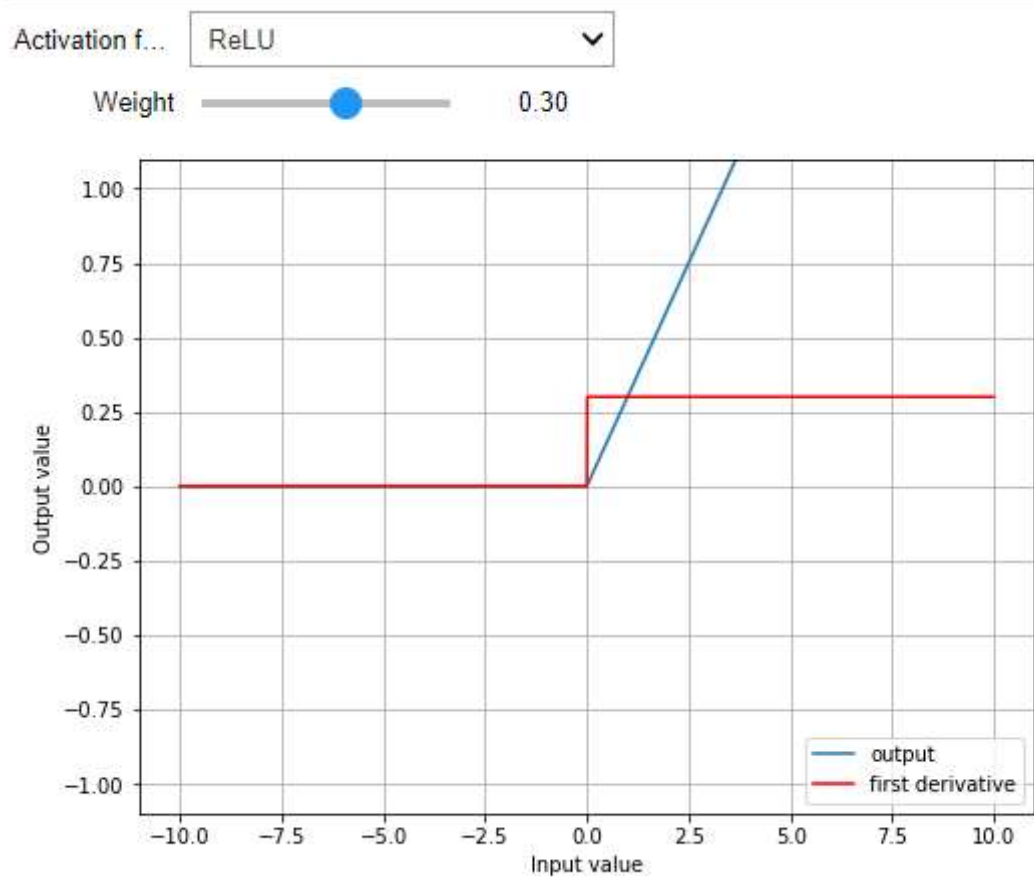


Abdalla Farid & Noah Graells

ReLU

In []:


```
def relu(neta):  
    output = np.array([0 if x < 0 else x for x in neta])  
    d_output = np.array([0 if x < 0 else 1 for x in neta])  
    return (output, d_output)
```



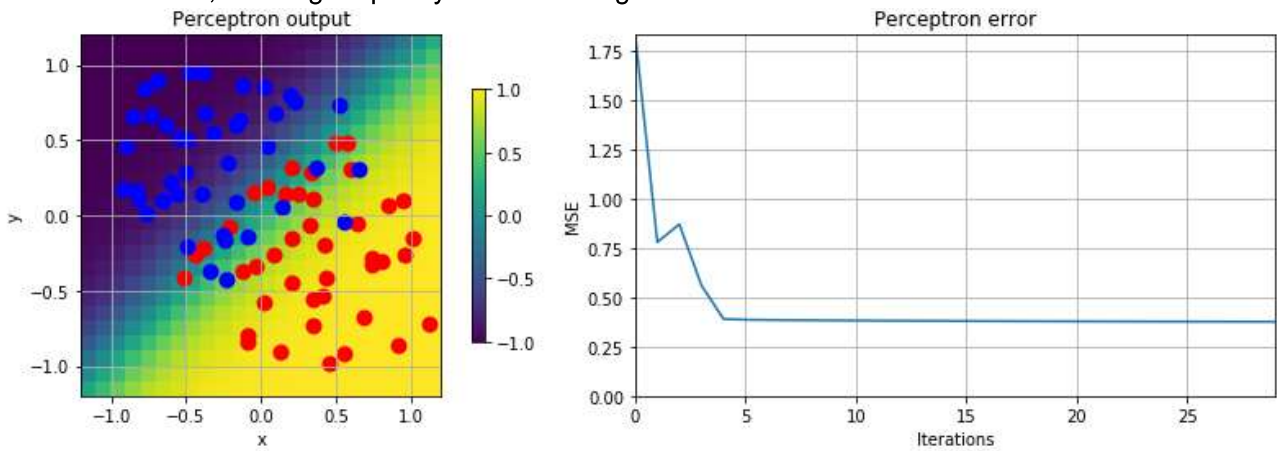
4_delta-rule

Answers

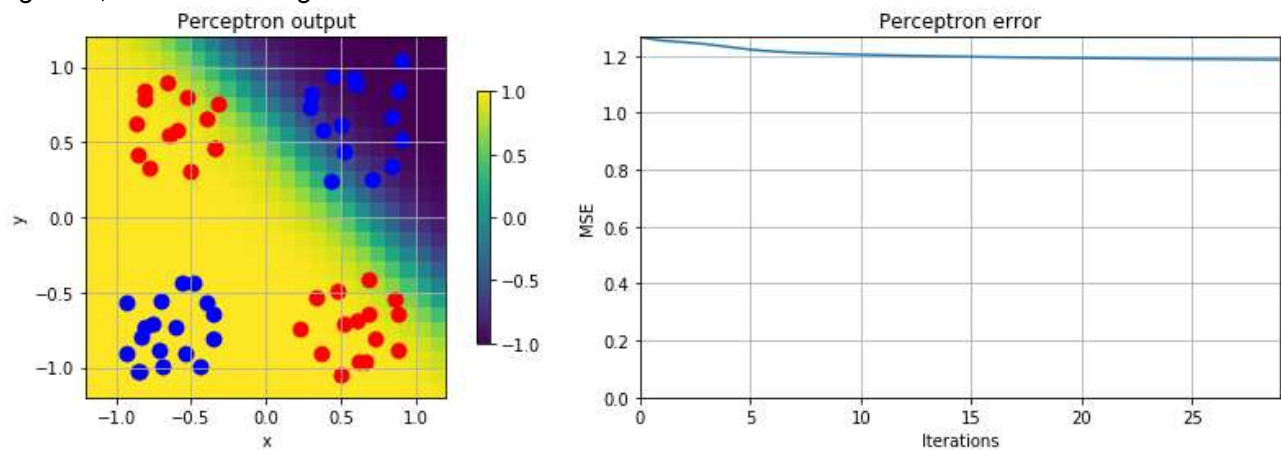
1. No oscillation, converges quickly and finds a good solution (low error)



2. Small oscillations, converges quickly. Error not as good as the first one but still ok.

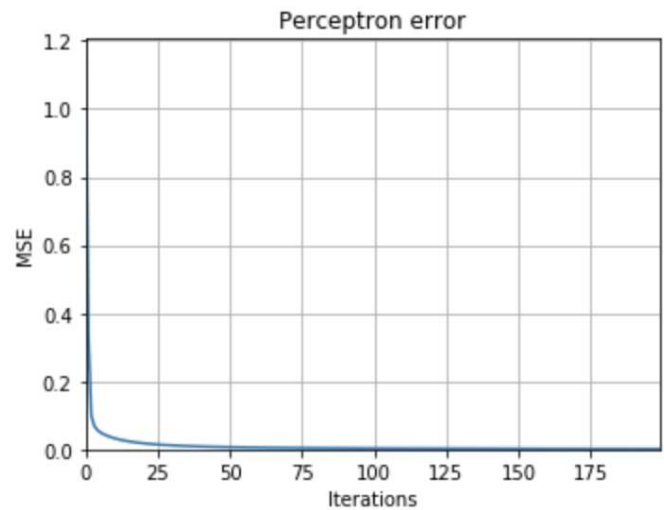



3. Big error, doesn't converge.

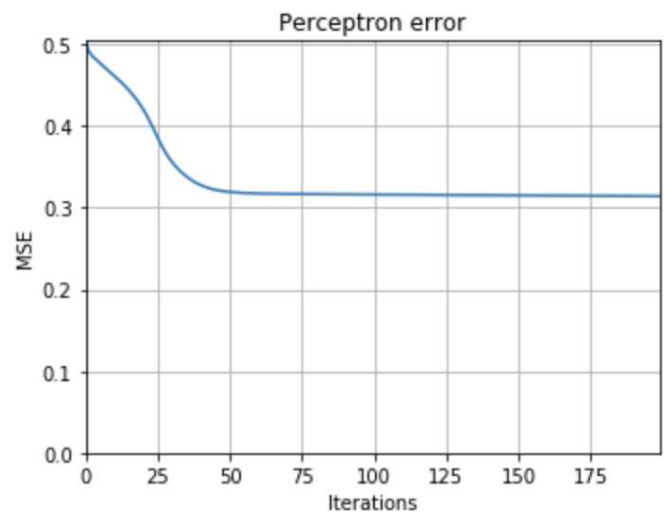



5_backpropagation


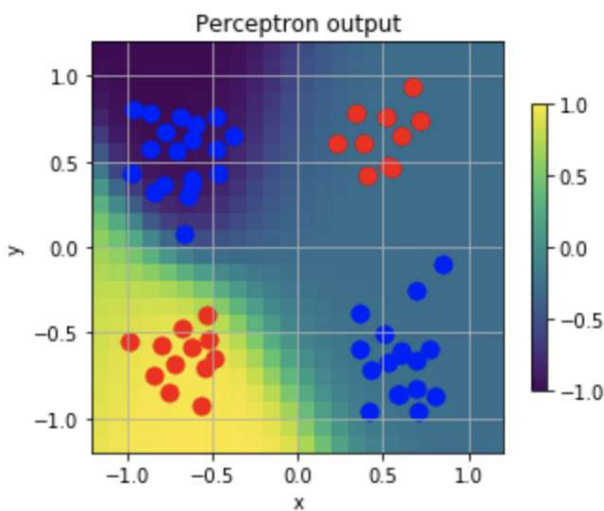
When classes are separable, we reach quickly the optimal solution with a really low error :



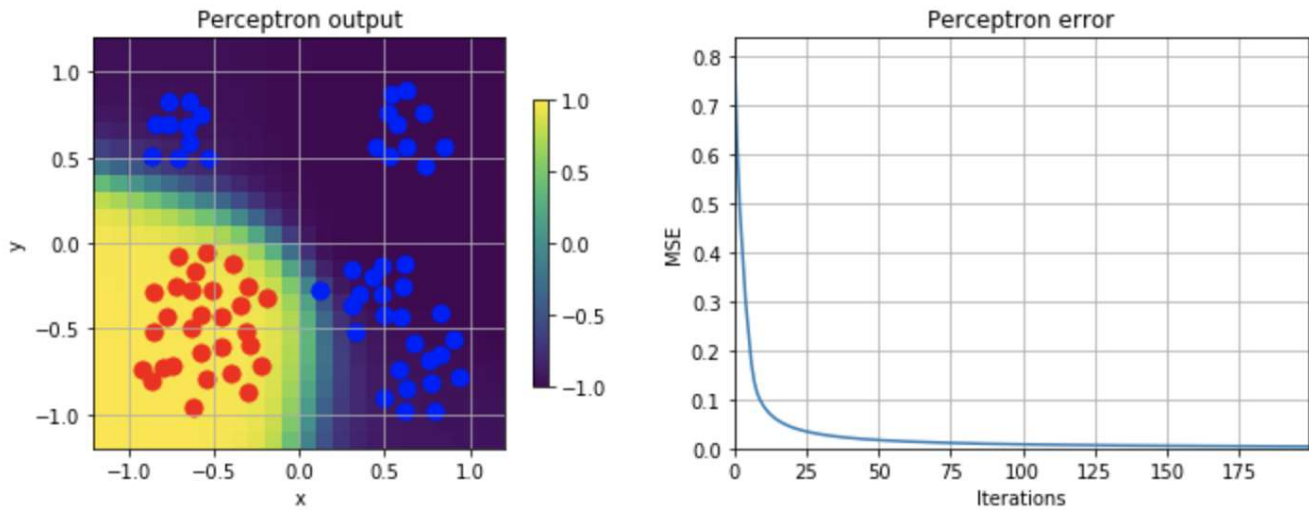
With overlapping classes, it converges quite quickly but the error is quite big since the classes aren't well defined. The convergence is smoother than without the backpropagation :



With non-separable classes, the error is big but not as big as the exercise 4. There is not really a local minima since there is an horizontal asymptote :



The result is similar to the first question since classes are separable. The convergences is not as quick though :



6_backpropagation_momentum

In [4]:

```

class MLP:
    """
    This code was adapted from:
    https://rolisz.ro/2013/04/18/neural-networks-in-python/
    """

    def __tanh(self, x):
        """Hyperbolic tangent function"""
        return np.tanh(x)

    def __tanh_deriv(self, a):
        """Hyperbolic tangent derivative"""
        return 1.0 - a**2

    def __logistic(self, x):
        """Sigmoidal function"""
        return 1.0 / (1.0 + np.exp(-x))

    def __logistic_derivative(self, a):
        """sigmoidal derivative"""
        return a * (1 - a)

    def __init__(self, layers, activation='tanh'):
        """
        :param layers: A list containing the number of units in each layer.
        Should be at least two values
        :param activation: The activation function to be used. Can be
        "logistic" or "tanh"
        """
        self.n_inputs = layers[0] # Number of inputs (first layer)
        self.n_outputs = layers[-1] # Number of outputs (last layer)
        self.layers = layers # Activation function

        if activation == 'logistic':
            self.activation = self.__logistic
            self.activation_deriv = self.__logistic_derivative
        elif activation == 'tanh':
            self.activation = self.__tanh
            self.activation_deriv = self.__tanh_deriv

        self.init_weights() # Initialize the weights of the MLP

    def init_weights(self):
        """
        This function creates the matrix of weights and initializes their values to small values
        """
        self.weights = [] # Start with an empty list
        for i in range(1, len(self.layers) - 1): # Iterates through the layers
            # np.random.random((M, N)) returns a MxN matrix
            # of random floats in

```

```

[0.0, 1.0).
# (self.layers[i] + 1)
is number of neurons in layer i plus the bias unit
self.weights.append((2 * np.random.random((self.layers[i - 1] + 1, self.lay
ers[i] + 1)) - 1) * 0.25)
# delta_weights are ini
tialized to zero
# Append a last set of
weights connecting the output of the network
self.weights.append((2 * np.random.random((self.layers[i] + 1, self.layers[i +
1])) - 1) * 0.25)

def fit(self, data_train, data_test=None, learning_rate=0.1, momentum=0.5, epochs=1
00):
    '''
    Online Learning.
    :param data_train: A tuple (X, y) with input data and targets for training
    :param data_test: A tuple (X, y) with input data and targets for testing
    :param learning_rate: parameters defining the speed of Learning
    :param epochs: number of times the dataset is presented to the network for Lear
ning
    '''
    X = np.atleast_2d(data_train[0])
    temp = np.ones([X.shape[0], X.shape[1]+1])
    # Inputs for training
    # Append the bias unit
    to the input layer
    temp[:, 0:-1] = X
    X = temp
    # X contains now the in
puts plus a last column of ones (bias unit)
    y = np.array(data_train[1])
    error_train = np.zeros(epochs)
    # Targets for training
    # Initialize the array
    to store the error during training (epochs)
    if data_test is not None:
        # If the test data is p
rovided
        error_test = np.zeros(epochs)
        # Initialize the array
        to store the error during testing (epochs)
        out_test = np.zeros(data_test[1].shape)
        # Initialize the array
        to store the output during testing

    a = []
    # Create a list of arra
ys of activations
    for l in self.layers:
        a.append(np.zeros(1))
    # One array of zeros pe
r Layer

    for k in range(epochs):
        # Iterate through the e
pochs
        error_it = np.zeros(X.shape[0])
        # Initialize an array t
o store the errors during training (n examples)

        for it in range(X.shape[0]):
            # Iterate through the e
xamples in the training set
            i = np.random.randint(X.shape[0])
            # Select one random exa
mple
            a[0] = X[i]
            # The activation of the
first layer is the input values of the example

            # Feed-forward
            for l in range(len(self.weights)):
                # Iterate and compute t
he activation of each layer

```

```

        a[l+1] = self.activation(np.dot(a[l], self.weights[l])) # Apply the
activation function to the product input.weights

        error = a[-1] - y[i] # Compute the error: ou
tput - target
        error_it[it] = np.mean(error ** 2) # Store the error of th
is iteration (average of all the outputs)
        deltas = [error * self.activation_deriv(a[-1])] # Ponderate the error b
y the derivative = delta

# Back-propagation
# We need to begin at t
he layer previous to the last one (out->in)
        for l in range(len(a) - 2, 0, -1): # Append a delta for ea
ch layer
            deltas.append(deltas[-1].dot(self.weights[l].T) * self.activation_d
eriv(a[l]))
            deltas.reverse() # Reverse the list (in-
>out)

# Update
# Iterate through the l
ayers
        for i in range(len(self.weights)):
            previous_delta_weights = 0
            layer = np.atleast_2d(a[i]) # Activation
            delta = np.atleast_2d(deltas[i]) # Delta
            # Compute the weight ch
ange using the delta for this layer
            # and the change comput
ed for the previous example for this layer
            delta_weights = -learning_rate * layer.T.dot(delta) + momentum * pr
evious_delta_weights
            self.weights[i] += delta_weights # Update the weights
            previous_delta_weights = delta_weights

        error_train[k] = np.mean(error_it) # Compute the average o
f the error of all the examples
        if data_test is not None: # If a testing dataset
was provided
            error_test[k], _ = self.compute_MSE(data_test) # Compute the testing e
rror after iteration k

        if data_test is None: # If only a training da
ta was provided
            return error_train # Return the error duri
ng training
        else:
            return (error_train, error_test) # Otherwise, return bot
h training and testing error

    def predict(self, x):
        """
        Evaluates the network for a single observation
        """
        x = np.array(x)
        temp = np.ones(x.shape[0]+1)
        temp[0:-1] = x
        a = temp

```

```

for l in range(0, len(self.weights)):
    a = self.activation(np.dot(a, self.weights[l]))
return a

def compute_output(self, data):
    """
    Evaluates the network for a dataset with multiple observations
    """
    assert len(data.shape) == 2, 'data must be a 2-dimensional array'


    out = np.zeros((data.shape[0], self.n_outputs))
    for r in np.arange(data.shape[0]):
        out[r,:] = self.predict(data[r,:])
    return out

def compute_MSE(self, data_test):
    """
    Evaluates the network for a given dataset and
    computes the error between the target data provided
    and the output of the network
    """
    assert len(data_test[0].shape) == 2, 'data[0] must be a 2-dimensional array'


    out = self.compute_output(data_test[0])
    return (np.mean((data_test[1] - out) ** 2), out)

```

7_hold_out_validation



8_cross_validation

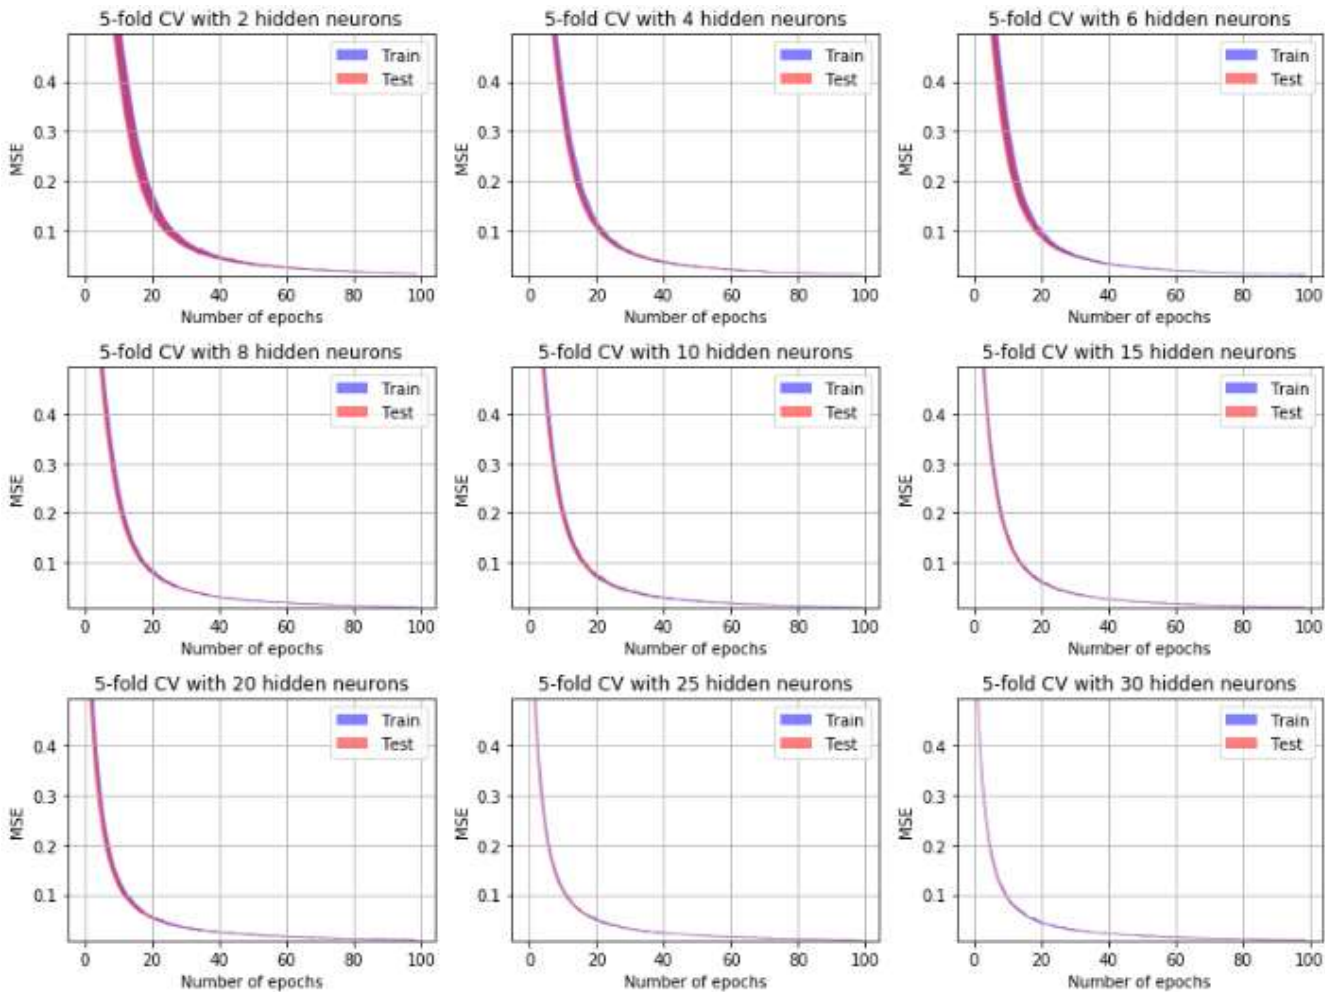



Difference between 7 and 8

We can see that with the cross validation there are less dispersions than with the hold-out validation. The median error is better for the hold out validation (overfitting).

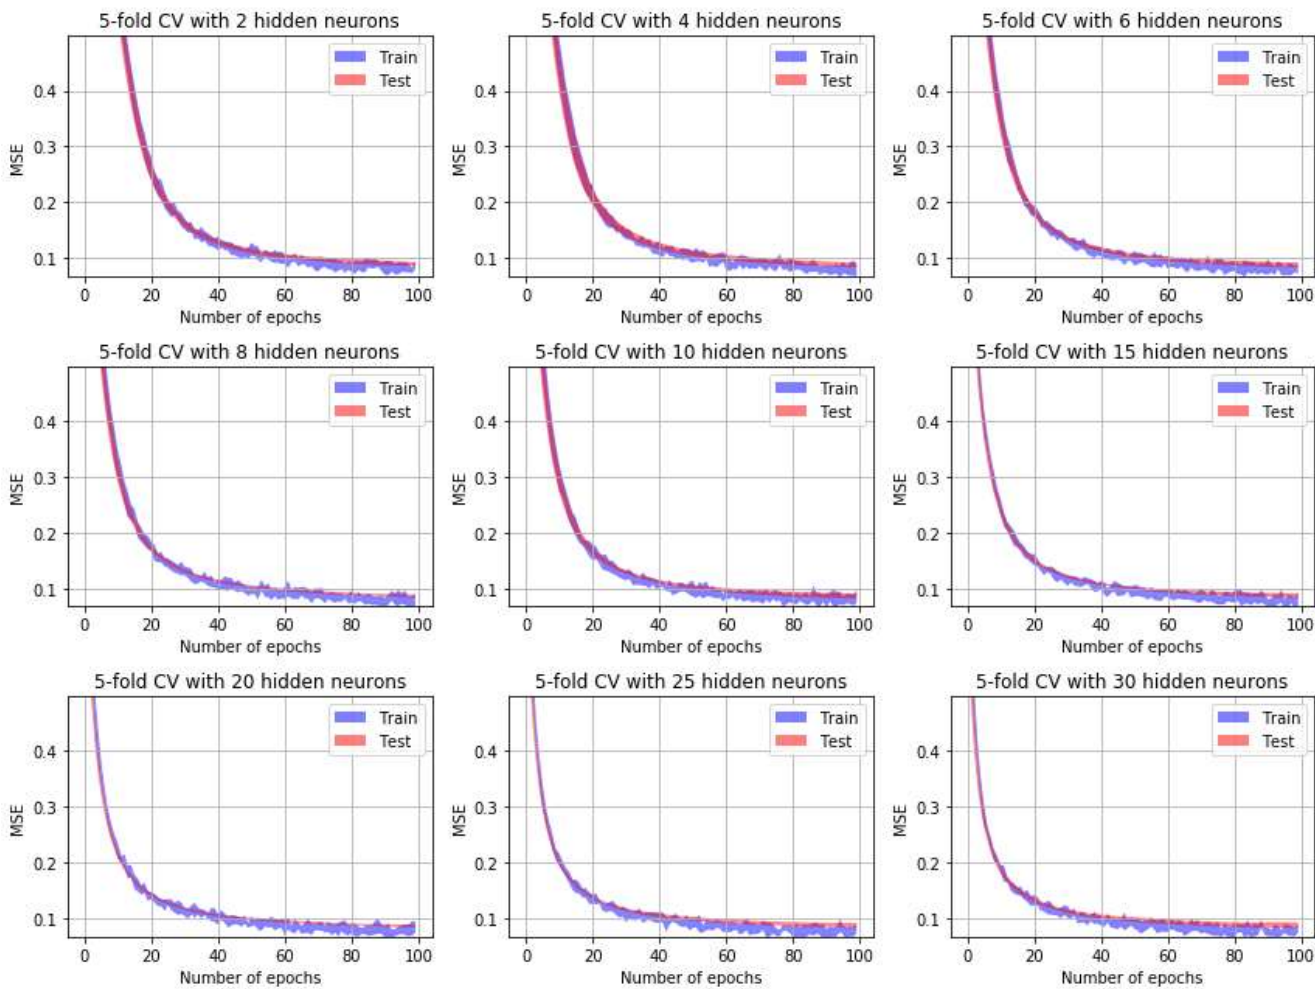
9_Model_Selection



Spread 0.3

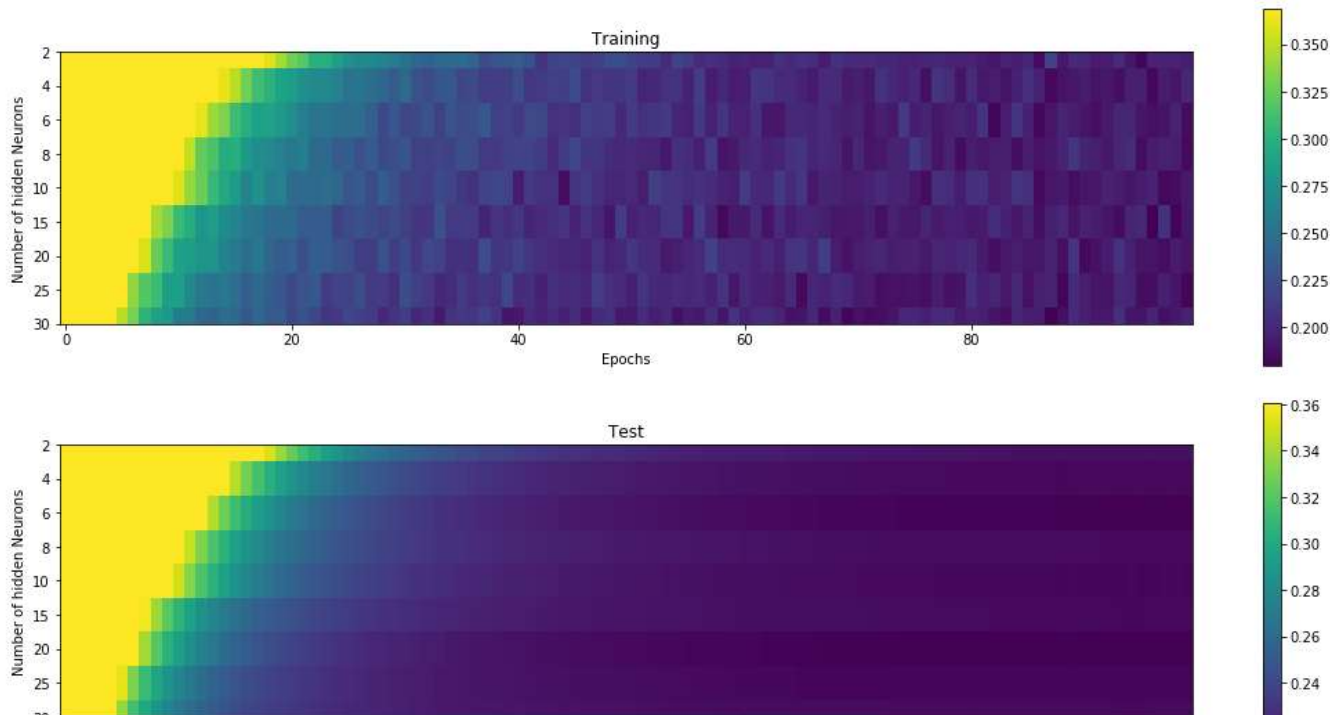




Spread 0.5







4 hidden neurons is a good tradeoff for convergence time (around 60 epochs) and computation time (since more neurons = more computation time). If we look at the thickness of the red line, 4 neurons seems better (2 neurons results in a thick line for a spread 0.7)

In []: